

High Level Design Document  
For  
Multimodal Education Creator  
**BACHELOR OF TECHNOLOGY in COMPUTER SCIENCE & ENGINEERING**  
BY

| Sr.no | Student Name   | Enrolment Number |
|-------|----------------|------------------|
| 1     | ABHIKET THAKUR | EN22CS301030     |
| 2     | ADITYA PALIWAL | EN22CS301055     |
| 3     | AARUSHI MITTAL | EN24CS3T10026    |
| 4     | AKASH KUSHWAHA | EN22CS301083     |

**Datagami Skill Based Course**



**MEDICAPS**  
UNIVERSITY

**Group: 04D1**

**Project No: GAI-04**

**Department of Computer Science & Engineering  
Faculty of Engineering  
MEDICAPS UNIVERSITY, INDORE- 453331  
Jan - May 2026**

## 1. Introduction

### 1.1 Scope of the Document

This document presents the High-Level Design (HLD) of the AI-Powered Educational Content Creator System. It provides a comprehensive overview of the system architecture, major components, data flow, technology stack, and deployment structure.

The purpose of this document is to:

- Define the architectural design of the system.
- Describe the interaction between major modules.
- Provide technical clarity for evaluators.
- Serve as a reference for development, maintenance, and future enhancements

### 1.2 Intended Audience

The primary audience includes:

- Project Evaluators and Academic Reviewers
- Development Team Members
- Future Maintainers and Enhancers
- Research-Oriented Audience

### 1.3 System overview

The AI-Based Educational Content Creator is an intelligent web-based system designed to generate structured and context-aware educational material using Natural Language Processing (NLP) and semantic vector retrieval techniques. The system leverages a Retrieval-Augmented Generation (RAG) architecture to enhance the relevance and accuracy of generated educational content by retrieving semantically similar knowledge from a pre-indexed vector database.

**Purpose of the System:** The primary objective of the system is to,

- Automate the generation of structured educational content
- Reduce manual effort in content creation
- Improve contextual relevance using semantic search
- Demonstrate practical implementation of NLP and vector-based retrieval

**Core Functionality:** The system performs the following high-level operations,

- Accepts a user-provided topic or prompt through a web interface
- Preprocesses and converts the input into a semantic embedding
- Performs similarity search on a precomputed vector database
- Retrieves relevant educational content chunks
- Constructs structured and formatted educational material

**Architectural Approach:** The system follows a layered modular architecture, consisting of,

- Presentation Layer (Web Interface)
- Application & NLP Processing Layer
- Vector Retrieval Layer
- Knowledge Storage Layer


**Design Philosophy:** The system is designed with,

- **Modularity** – Each component can be independently upgraded
- **Performance Optimization** – Precomputed embeddings reduce runtime cost
- **Scalability Readiness** – Architecture supports migration to distributed vector databases
- **Minimal Data Retention** – No permanent storage of user data


## 2. System Design

### 2.1 Application Design

The application follows a layered modular architecture implementing a Retrieval-Augmented Generation (RAG) workflow. It separates user interaction, processing logic, and data storage to ensure maintainability and scalability.



## 2.2 Process Flow



The system converts the user query into embeddings, retrieves semantically relevant content from the vector database, and generates structured educational material using a RAG-based workflow.

- **User Input:** User enters a topic or prompt through the web interface.
- **Input Validation & Preprocessing:** Text is cleaned, normalized, and prepared for embedding generation.
- **Embedding Generation:** The processed query is converted into a semantic vector representation.
- **Vector Similarity Search:** Query embedding is compared with stored content embeddings using cosine similarity to retrieve Top-K relevant chunks.
- **Context Assembly:** Retrieved content chunks are combined to form contextual knowledge.
- **Content Generation:** Structured educational material is generated using retrieved context.
- **Response Display:** Final formatted output is returned and displayed to the user.

## 2.3 Information Flow

Information flows through the system in a structured, sequential manner — transforming raw user input into semantically enriched educational content.

### Information Movement Across Layers:

#### User Layer → Presentation Layer

- User submits topic or prompt
- Request is captured by the web interface

#### Presentation Layer → Processing Layer

- Input is validated and preprocessed
- Cleaned text is forwarded for embedding generation

### Processing Layer → Data Layer


- Query is converted into embedding vector
- Vector is compared with stored content embeddings
- Top-K relevant content chunks are retrieved

### Data Layer → Processing Layer

- Retrieved contextual information is assembled
- Structured educational content is generated

### Processing Layer → Presentation Layer


- Final formatted output is returned
- Content is displayed to the user



## 2.4 Components Design

The system is designed using modular components, where each component has a clearly defined responsibility. This ensures maintainability, scalability, and independent upgrades.

- **Web Interface Component:** Captures user topic or prompt and Sends request to backend. Displays structured educational output and Maintains request-response lifecycle.
- **Input Processing Component:** Validates and sanitizes user input and Performs text cleaning and normalization. Prepares input for embedding generation.
- **Embedding Component:** Converts processed text into semantic vector and Enables similarity-based retrieval. Represents contextual meaning in high-dimensional space.
- **Retrieval Component:** Performs cosine similarity computation and Filters based on similarity score. Retrieves Top-K relevant content chunks.
- **Content Construction Component:** Combines retrieved context and Formats headings and explanations. Structures content into organized educational format.
- **Vector Database Component:** Stores precomputed embeddings and supports fast semantic search and stores content chunks and metadata.



## 2.5 Key Design Considerations

**1. Modularity:** The system is divided into independent components (UI, NLP, Retrieval, Vector DB) to ensure maintainability and easy upgrades.

### 2. Performance Optimization

- Precomputed embeddings reduce runtime computation.
- In-memory vector loading minimizes latency.

- Top-K retrieval limits unnecessary processing.

### 3. Scalability

- Architecture supports migration to distributed vector databases (e.g., FAISS).
- Can integrate caching and load balancing for high traffic.
- Suitable for containerized cloud deployment.

### 4. Accuracy & Relevance

- Uses semantic embeddings instead of keyword matching.
- Implements Retrieval-Augmented Generation (RAG) for context-aware output.
- Reduces hallucination by grounding responses in stored knowledge.

### 5. Security & Data Privacy

- Input validation and sanitization.
- Controlled access to serialized vector database.
- No permanent storage of user prompts.

### 6. Maintainability

- Clear separation of concerns.
- Replaceable embedding model or storage system.
- Structured codebase with defined module responsibilities.

### 7. Stateless Processing

- Each query processed independently.
- No cross-session data sharing.
- Simplifies scaling and deployment.

## 2.6 API Catalogue

The system exposes internal application APIs to handle user requests, process queries, perform semantic retrieval, and generate structured educational content.

- Generate Educational Content:** Accepts a user topic or prompt and returns structured educational material using RAG-based retrieval.
  - Endpoint: ***POST /generate-content***
  - Request Body (JSON): `{"topic": "Machine Learning"}`
  - Response (JSON): `{"status": "success", "generated_content": "Structured educational explanation..."}`
- Health Check API:** Verifies that the application and vector database are loaded and running.

- Endpoint: **GET /health**
- Response: `{"status": "active"}`


**III. Cache Status API:** Returns current cache usage details for performance monitoring.

- Endpoint: **GET /cache-status**

### 3. Data Design

#### 3.1 Data Model

The data model is designed to support semantic retrieval and structured content generation. It separates user input, embeddings, stored knowledge, and generated output for modular and scalable processing.



#### Data Relationships:

- Each UserQuery generates one QueryEmbedding
- Each ContentChunk has one ContentEmbedding
- Query embeddings perform similarity search on content embeddings
- Retrieved content produces GeneratedContent

## 3.2 Data Access Mechanism

The system uses a vector-based data access mechanism to retrieve semantically relevant educational content. Instead of traditional SQL queries, the system performs embedding-based similarity search.

### Data Storage Strategy

- Educational content is segmented into content chunks
- Each chunk has a precomputed embedding
- Stored in a serialized vector database (education\_vector\_db.pkl)
- Loaded into memory during application startup

### Data Retrieval Process

- User query is converted into embedding
- Cosine similarity is computed against stored embeddings
- Top-K most relevant content chunks are retrieved
- Retrieved data is passed to content generation module

### Access Characteristics

- Read-heavy design (retrieval-focused)
- In-memory vector comparison for low latency
- No direct modification of stored embeddings at runtime
- Stateless query-based access

### Performance Optimization

- Precomputed embeddings reduce real-time computation
- Top-K filtering minimizes unnecessary data processing
- Optional caching for repeated queries

### Security Considerations

- Controlled loading of serialized .pkl file
- Restricted file access permissions
- No permanent storage of user queries

## 3.3 Data Retention Policies

### Data Lifecycle Process

#### Step 1 – Data Collection

- Educational materials are uploaded by administrator.
- User inputs are captured during active sessions.

#### Step 2 – Processing

- Content is chunked and converted into embeddings.
- Embeddings are stored in vector database.

### Step 3 – Active Usage

- Data is accessed for similarity search and response generation.
- Query logs may be stored for analytics.

### Step 4 – Archival / Deletion

- Old logs are automatically deleted after retention period.
- Dataset updates trigger embedding regeneration.
- Session-based content is cleared automatically.

### Retention Rules:

- No sensitive personal data is permanently stored.
- User session data is cleared after logout or timeout.
- Logs older than 30–90 days are automatically purged.
- If dataset is replaced, old embeddings are deleted to prevent redundancy.
- Backup retention period: 7–14 days.

### Security & Compliance Measures:

- Role-based access control for administrators.
- Encrypted storage for logs (if deployed in production).
- Regular backup and integrity checks.
- Compliance with general data protection principles

### Data Deletion Policy

Data is deleted under the following conditions:

- Manual administrator request
- Dataset replacement
- System reset
- Expired retention period
- User account deletion

Deletion includes:

- Removing embeddings
- Clearing associated logs
- Cleaning temporary files

## 3.4 Data Migration

Data migration refers to the structured process of transferring educational datasets, embeddings, and system configurations from one environment to another

### The system ensures safe migration of:

- Educational content datasets
- Vector embeddings
- Model configurations
- Application metadata

## Migration Scenarios

- Initial dataset deployment
- Dataset updates (adding new study material)
- Migration between environments (local → cloud)
- Vector database upgrade or format change
- Backup restoration

## Migration Process

### Data Export

- Export existing dataset files
- Backup vector database (education\_vector\_db.pkl / FAISS index)
- Export configuration settings

### Data Transfer

- Secure transfer using encrypted channels (SCP / SFTP / Cloud Storage)
- Maintain file integrity using checksum validation

### Data Import

- Load dataset into new environment
- Regenerate embeddings if model version changes
- Validate vector database compatibility

### Validation

- Run similarity search test queries
- Validate content retrieval accuracy
- Perform API health check

## 4. Interfaces

The system exposes structured interfaces to enable interaction between users, application modules, and external services. Interfaces are designed using REST principles and modular internal communication.

**User Interface:** Web-based interface (HTML / CSS / JS / Frontend Framework)

- Accept user topic input
- Display generated educational content
- Show error or system status messages
- Sends HTTP requests to backend APIs
- Receives JSON responses

**Application Programming Interface (API):** RESTful Web API, JSON request/response

- Accept user topic
- Trigger embedding generation
- Perform retrieval
- Return structured output

### Core Endpoints:

- *POST /generate-content*
- *GET /health*
- *GET /cache-status* (optional)

## Data Interface

### Vector Database Interface:

- Loads. pkl or FAISS index file
- Performs cosine similarity search
- Returns Top-K relevant content chunks

### File System Interface:

- Loads educational dataset
- Accesses stored embeddings
- Reads configuration files

## External Service Interfaces

- Embedding Model API (e.g., local transformer or cloud model)
- Large Language Model API (for content generation)
- Optional Cloud Storage Service

## Interface Characteristics

- Stateless REST communication
- JSON-based data exchange
- Secure HTTP (HTTPS) recommended
- Modular and loosely coupled design
- Easily extensible

## **5. State and Session Management**

- **Stateless Backend Design:** Each user request (topic input) is processed independently without storing persistent session state on the server.
- **Session-Level Context Handling:** Temporary session memory can maintain user interaction history during active usage (if enabled).
- **In-Memory Runtime State:** Vector database and cache are loaded into memory during application runtime for faster processing.
- **Session Isolation:** User inputs are handled separately to prevent cross-session data leakage.

- **Future Enhancement:** Can integrate JWT-based authentication and server-side session storage for multi-user production deployment

## **6. Caching**

### **Caching Strategy**

- **Purpose of Caching:** Reduces repeated computation and improves response time for frequently requested topics.
- **Embedding Cache:** Stores previously generated query embeddings to avoid redundant vector computation.
- **Retrieval Result Cache:** Saves Top-K similarity search results for commonly searched educational topics.
- **In-Memory Caching Mechanism:** Fast access using in-memory storage for low-latency responses.
- **Scalable Extension:** Can be extended using distributed caching systems (e.g., Redis) in cloud deployment for high traffic handling.

### **Caching Mechanism in Educational Content Creator**

- **Query Embedding Cache:** When a user enters a topic, its embedding is cached to avoid regenerating embeddings for repeated or similar queries.
- **Vector Retrieval Cache:** Frequently searched topics store their Top-K retrieved content chunks, reducing repeated similarity computations.
- **Generated Content Cache:** Final structured educational content for common topics is temporarily cached to improve response time.
- **In-Memory Storage:** Cache operates in-memory for low-latency access during runtime.
- **Future Scalability:** Can be extended using distributed caching (e.g., Redis) for cloud-based, high-traffic deployment.

## **7. Non-Functional Requirements**

### **7.1 Security Aspects**

Security is an essential consideration in the design of the AI-Based Educational Content Creator. Although the system operates primarily as a content generation and retrieval platform, it handles user input, backend processing, and vector-based knowledge storage. Therefore, appropriate security controls are implemented to ensure system integrity, data protection, and controlled access.

- **Input Validation and Sanitization:** Since the system accepts user-generated prompts through the web interface:

- All input data is validated before processing.
- Special characters and malicious patterns are sanitized.
- Protection mechanisms are applied to prevent:
  - Injection attacks
  - Script-based manipulation
  - Unauthorized command execution

This ensures the NLP pipeline processes only clean and structured input.

- **Secure Handling of Vector Database (.pkl):** The system uses a serialized vector database file (education\_vector\_db.pkl) to store embeddings and content metadata. Security measures include:

- Restricted file access permissions Controlled loading of serialized object
- Prevention of arbitrary file execution
- Secure storage within the application directory

- **Application-Level Security:** The backend architecture ensures:

- Separation of presentation layer and processing layer
- Controlled interaction between web interface and NLP module
- Encapsulation of embedding and retrieval logic

- **Authentication & Access Control (Future Enhancement):** For scalable or production deployment, the system can integrate:

- User authentication (JWT-based session handling)
- Role-based access control (RBAC)
- API-level authorization tokens

- **Data Privacy Considerations:** This ensures minimal data retention risk.

- No personally identifiable information (PII) is permanently stored.
- User prompts are processed in-memory.
- No external sharing of user data occurs.

- **Secure Deployment Practices:** For production-level deployment, recommended measures include:

- HTTPS encryption for secure communication
- Reverse proxy configuration (e.g., Nginx)
- Containerization using Docker for isolation
- Secure environment variable management
- Firewall configuration on cloud servers

- **Model & Retrieval Security:** To prevent misuse of the AI pipeline:

- Input size limits can be enforced
- Rate limiting can be implemented
- Logging and monitoring mechanisms can be added
- Retrieval responses can be filtered for safe output formatting

- **Dependency and Environment Security:**

- Virtual environment (.venv) isolates dependencies
- requirements.txt ensures controlled package installation
- Regular updates prevent known vulnerabilities

## 7.2 Performance Aspects

- **Precomputed Embeddings:** Educational content embeddings are generated and stored in advance, reducing real-time computation overhead.
- **Efficient Similarity Search:** Semantic retrieval using cosine similarity enables fast Top-K relevant content fetching.
- **In-Memory Processing:** Vector database is loaded into memory for low-latency retrieval.
- **Modular Processing Pipeline:** Separation of NLP, retrieval, and generation improves execution efficiency.
- **Scalability Ready:** Can be optimized further using FAISS, caching mechanisms, parallel processing, and cloud deployment.

## 8. References

- Lewis, P., et al. (2020). *Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks*. Advances in Neural Information Processing Systems (NeurIPS).
- Reimers, N., & Gurevych, I. (2019). *Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks*. Proceedings of EMNLP-IJCNLP.
- Johnson, J., Douze, M., & Jégou, H. (2017). *FAISS: A library for efficient similarity search and clustering of dense vectors*. Facebook AI Research.
- FastAPI Documentation. <https://fastapi.tiangolo.com/>
- Hugging Face Transformers Documentation. <https://huggingface.co/docs/transformers>
- scikit-learn Documentation. <https://scikit-learn.org/>
- NumPy Documentation. <https://numpy.org/>
- Fielding, R. T. (2000). *Architectural Styles and the Design of Network-based Software Architectures (REST)*. Doctoral Dissertation, University of California, Irvine.
- Designing Data-Intensive Applications, Kleppmann, M. (2017), O'Reilly Media.
- OpenAI API Documentation (for LLM integration concepts). <https://platform.openai.com/docs>